

Machine Learning.

Angel Chaico

yo

Facultad de ciencias, UNI¹

23 de julio de 2022

¹yo

Índice general

1. aora	9
I Statistical and probabiliti (EDA).	11
1.1. Estimates and explorind data.	13
1.2. Distributions.	13
1.3. Statistical experiment.	13
II Resamplin(train and test set).	15
2. Sampli:stratied	17
3. Imbalance data.	19
3.1. Undersampling	19
3.2. Oversamplig	19
3.2.1. Bootstraping	19
3.2.2. data generation: SMOTE algorithm	19
3.3. Weighting.	19
3.4. Cost-based classification.	19

3.5. bibliografi	19
III Supervised (model)	21
4. Linear regression.	23
5. Logistic regression.	25
5.1. Regularization.	25
5.2. Código.	25
6. Desicion tree.	27
6.1. regression	27
6.2. Classification	28
7. Random Forest.	29
7.1. Regularization.	29
7.2. Código.	29
8. Bagging	31
9. Boosting.	33
10. Bayesian additive regression trees.	35
11. Supert Vector Machine.	37
11.1. Regularization.	37
11.2. Código.	37

<i>ÍNDICE GENERAL</i>	5
IV Unsupervised (model)	39
12.K-means	41
13.Hierarchical Clustering.	43
V Deep learning (model)	45
14.Neural Networks	47
14.1. Activation function.	47
14.2. Convulation layer.	47
14.3. Poling layer.	47
15.Regularization.	49
15.1. The Vanishing/Exploring gradients	49
15.2. Reusing Petrained Layer	52
15.3. Faster Optimizationers.	52
15.4. Overfitting.	52
15.5. Summary	54
15.6. bibliografy	54
VI Optimization (opt model).	55
16.First order methods.	57
16.1. Adagrad.	57
16.2. Adam.	57

17.Second order methods.	59
17.1. Newton's method.	59
17.2. Secant method.	59
18.Direct Methods.	61
18.1. Nelder-mead simplex method.	61
19.Stochastic methods.	63
19.1. Noise descent.	63
19.2. Cross-Entropy method.	63
VII Regularization (overfitting).	65
VIII Metrics (evaluation)	67
20.Regression	69
21.Classification	71
21.1. Matrix confusion	71
21.2. ROC curve.	72
21.3. PR curve	74
21.4. Gain chart and lift chart.	74
21.5. K-S Chart	76
21.6. bibliografia.	76

<i>ÍNDICE GENERAL</i>	7
IX Python.	77
22.sklearn.preprocessing	79
23.sklearn.pipeline	81
24.sklearn.ensemble	83
24.1. sklearn.ensemble.RandomForestClassifier	83
25.sklearn.model_selection	85
26.sklearn.metrics	87
27.tf.keras.model	89
28.ft.keras.layer	91
X R.	93
29.ggplot	95
30.dplyr	97
31.tidyr	99

XI	SQL.	101
XII	Julia.	103
XIII	C++.	105

Capítulo 1



aora

random forest, neural network, regresion logisticy hacer bibliografi por capitulo.

En loc codigos iran ejemplos graficos de cada modelo aplicado.

Parte I

Statistical and probabiliti (EDA).



1.1. Estimates and explorind data.

1.2. Distributions.

1.3. Statistical experiment.

Parte II

Resamplin(train and test set).

Capítulo 2

Sampli:stratied

Capítulo 3

Imbalance data.

3.1. Undersampling

3.2. Oversampling

3.2.1. Bootstrapping

3.2.2. data generation: SMOTE algorithm

3.3. Weighting.

3.4. Cost-based classification.

3.5. bibliografía

revisar
practical statistical for data science, by peter bruce.

Parte III

Supervised (model)

Capítulo 4

Linear regression.

Capítulo 5

Logistic regression.

5.1. Regularization.

5.2. Código.

Capítulo 6

Decision tree.

6.1. regression

:

- We divide the predictor space- that is, the set of possible values for X_1, X_2, \dots, X_p - into J distinct and non-overlapping regions, R_1, R_2, \dots, R_J .
- For every observation that falls into the region R_j , we make the same prediction, which is simply the mean of the response values for the training observations in R_j .

The goal es to find R_1, \dots, R_J that minimize the RSS given by:

$$\sum_{j=1}^J \sum_{i \in R_j} (y_i - \hat{y}_{R_j})^2 \quad (6.1)$$

Recursive binary splitting, we seek the value of j and s that minimize the equation

$$\sum_{i: x_i \in R_1(j,s)} (y_i - \hat{y}_{R_1})^2 + \sum_{i: x_i \in R_2(j,s)} (y_i - \hat{y}_{R_2})^2 \quad (6.2)$$

Cost complexity pruning (weakest link pruning). For each value of α there corresponds a subtree $T \in T_0$, T_0 is when $\alpha = 0$, such that is as small as possible. The $|T|$ indicate the number of terminal nodes of the tree T .

$$\sum_{m=1}^{|T|} \sum_{i: x_i \in R_m} (y_i - \hat{y}_{R_m})^2 + \alpha |T| \quad (6.3)$$

Algorithm 8.1 *Building a Regression Tree*

1. Use recursive binary splitting to grow a large tree on the training data, stopping only when each terminal node has fewer than some minimum number of observations.
 2. Apply cost complexity pruning to the large tree in order to obtain a sequence of best subtrees, as a function of α .
 3. Use K-fold cross-validation to choose α . That is, divide the training observations into K folds. For each $k = 1, \dots, K$:
 - (a) Repeat Steps 1 and 2 on all but the k th fold of the training data.
 - (b) Evaluate the mean squared prediction error on the data in the left-out k th fold, as a function of α .Average the results for each value of α , and pick α to minimize the average error.
 4. Return the subtree from Step 2 that corresponds to the chosen value of α .
-

Figura 6.1: Although CV error is computed as a function of α , it is convenient to display the result as a function of $|T|$.

6.2. Classification

Tree vs linear

Capítulo 7

Random Forest.

When building decision trees, each time a split in a tree is considered, a random sample of m predictors is chosen as split candidates from the full set of p predictors (decorrelating trees).

7.1. Regularization.

- As with bagging, random forests will not overfit if we increase B trees, so in practice we use a value of B sufficiently large for the error rate to have settled down.

Metod	default	observation
random sample of m predictors	$m = \sqrt{p}$	
bootstrapped (max_sample)		
Out-of-bag (OOB)		
criterion	gini	
max depth		

7.2. Código.

Importance attributes

feature_importances_	
n_outputs_	

Observationa

- Para cada X_train diferente se obtiene features importances diferentes en el modelo resultante, como determinar las features importances si el modelo cambia al cambiar los X_train?

Capítulo 8

Bagging

Capítulo 9

Boosting.

Capítulo 10

Bayesian additive regression trees.

Capítulo 11

Supert Vector Machine.

11.1. Regularization.

11.2. Código.

Parte IV

Unsupervised (model)

Capítulo 12

K-means

Capítulo 13

Hierarchical Clustering.

Parte V

Deep learning (model)

Capítulo 14

Neural Networks

14.1. Activation function.

Sigmoid

ReLU

ELU

$$ELU_{\alpha} = \begin{cases} \alpha(e^z - 1) & , z < 0 \\ z & , z \geq 0 \end{cases} \quad (14.1)$$

14.2. Convulation layer.

14.3. Poling layer.

Capítulo 15

Regularization.

Problems:

- Vanishing gradient or exploding gradients.
- Not have enough training data.
- Slow.
- Overfitting.

15.1. The Vanishing/Exploring gradients

Glorot and He initialization

Glorot and Yoshua show that with logistic sigmoid function activation and weight initialization technique, the variance of the output of each layer is much greater than the variance of its inputs.

Glorot show that the connection weights of each layer must be initialized randomly as¹:

$$\text{Normal distribution, } \mu = 0, \sigma^2 = \frac{1}{fan_{avg}} \quad (15.1)$$

$$\text{Uniform distribution between } -r, r \text{ with } r = \sqrt{\frac{1}{fan_{avg}}} \quad (15.2)$$

$$fan_{avg} = \frac{fan_{int} + fan_{out}}{2} \quad (15.3)$$

¹fan_in is the number of input and fan_out the number of neuron of the layer

initialization	activation function	$\sigma^2(Normal)$
Glorot	None, tanh, logistic, softmax	$\frac{1}{fan_{avg}}$
He	ReLU and variant	$\frac{2}{fan_{int}}$
Lecun	SELU	$\frac{1}{fan_{int}}$

Cuadro 15.1: Initialization parameters for each type of activation function; $r = \sqrt{3\sigma^2}$.

```
keras.layers.Dense(10,, activation="relu", kernel_initializer="he_normal")
```

```
% fan_in->fan_avg
he_avg_init=keras.initializers.VarianceScaling(scale=2,mode="fan_avg")
keras.layers.Dense(10, activation="sigmoid", kernel_initializer=he_avg_init)
```

Nonsaturating Activation Function

Glorot and Bengio show that the problems with unstable gradients were in part due to a poor choice of activation function.

- **ReLU**, problem dying ReLU.
- **leaky ReLU**, $LeakyReLU_{\alpha}(z) = \max(\alpha z, z)$, usual fixed $\alpha = 0,01$, from one paper, it conclusions was that the leaky variants always outperform the estric ReLU, $\alpha = 0,2$ seemed to resutl in better performance.
- **RReLU**,randomize leaky ReLU, α is random during the training and fixed to an average value during testing; seems to act regularizer reducing the overfitting of the training set.
- **PReLU**, parametric leaky ReLU, α to be learned during training; was reported to strongly outperform ReLU on large image datasets, but on smaller it runs the risk of overfitting the training set.
- **ELU** exponential linear unit, better that all varians ReLU. ELU network will be sslower than a ReLU network.
- **SELU**, Scaled ELU, in stack of dense layer result self-normalize: the output of each layer will tend to preserve a mean of 0 and standar desviation of 1 during trainig, which solve the grading problem. condiction:input staandarized 0,1; initialized with LeCu; architecture must be sequential,all layer must be dense

```

keras.layers.Dense(10, kernel_initializer="he_normal")
keras.layers.LeakyReLU(alpha=0.2)\PReLU()
%
keras.layers.Dense(10, activation="selu", kernel_initializer="lecun_normal")

```

Bach Normalization

The before not guarantee that gradient come back during trainig. Sergey and Christian proposed BN, consist of adding an operation in the model just before of after the activation fuction of each hidden layer. The algorithm needs to estimate each input's mean and standar desviation (does over mini-bach)

$$\mu_B = \frac{1}{m_B} \sum_{i=1}^{m_B} x^{(i)} \quad (15.4)$$

$$\sigma_B^2 = \frac{1}{m_B} \sum_{i=1}^{m_B} (x^{(i)} - \mu_B)^2 \quad (15.5)$$

$$\hat{x}^{(i)} = \frac{x^{(i)} - \mu_B}{\sqrt{\sigma_B^2 + \epsilon}} \quad (15.6)$$

$$z^{(i)} = \gamma \cdot \hat{x}^{(i)} + \beta \quad (15.7)$$

γ (outputscaler vector), β (output offset vector) are learned through regular backpropagation, and μ (the final input mean vector) and σ (the final input standar desviation vector) are stimated using an exponential moving average (estimated during training but used after).

The vanishing gradients problem was strongly reduced; BN acts like a regularizer; there is a runtime penalty (evoiding updating the previos layer's weight and biases).

The aoutor of the BN paper arqued in favor of adding the BN layers before the activation function, rather than after. There is some debate about this, as which is preferable seems to depend on the task. Since BN layer include one offset parameter per input, you can remove the bias term from the previos layer.

momentum this parameter is used by the BN layer when it updates the exponential moving averages.

BN has become one of the most-used layers in deep neural networks.²

```
model = keras.models.Sequatial([
```

²Now fixed-updated by Hogyi Zhang is view recently

```

keras.layers.Flatten(input_shape=[28,28]),
keras.layers.BatchNormalization(),
keras.layers.Dense(300, kernel_initializer="he_normal", use_bias=False),
keras.layers.BatchNormalization(),
keras.layers.Activation("elu")
keras.layers.Dense(100, kernel_initializer="he_normal", use_bias=False),
keras.layers.BatchNormalization(),
keras.layers.Activation("elu"),
keras.layers.Dense(10, activation="softmax")
])

```

Gradient Clipping

```

optimizer= keras.optimizers.SGD(clipvalue=1.0)
model.compile(loss="mse", optimizer=optimizer)

```

15.2. Reusing Pretrained Layer

15.3. Faster Optimizers.

15.4. Overfitting.

Even though Batch Normalization was designed to solve the unstable gradients problems, it also acts like a pretty good regularizer.

L1,L2 regularization

```

layer = keras.layers.Dense(100, activation="elu", kernel_initializer="he_normal",
kernel_regularizer=keras.regularizers.l2(0.01))

```

```

from functools import partial
RegularizedDense = partial(
keras.layers.Dense,
activation="elu",
kernel_initializer="he_normal",

```

```

kernel_regularizer=keras.regularizers.l2(0.01)
)
model = keras.models.Sequential([
keras.layers.Flatten(input_shape=[28,28]),
RegularizedDense(300),
RegularizedDense(100),
RegularizedDense(10,activation="softmax",
kernel_initializer="glorot_uniform")
])

```

Dropout

At every training step, every neuron(including the input neurons, but always excluding the output neurons) has a probability p of being temporarily "dropped out", meaning it will be entirely ignored during this training step, but it may be active during the next step. The hyperparameter p is called the dropout rate, and it is typically set between %10 – %50. The state-of-the-art neural networks get 1 – 2% accuracy boost simply by adding dropout. Also, we need to multiply each input connection weight by the keep probability (1- p) after training. Alternatively, we can divide each neuron's output by the keep probability during training(don't equivalent but they work equally well).

```

model = keras.model.Sequential([
keras.layers.Flatten(input_shape=[28,28]),
keras.layers.Dropout(rate=0.2),
keras.layers.Dense(300, activation="elu",kernel_initializer="he_normal"),
keras.layers.Dropout(rate=0.2),
keras.layers.Dense(100, activation="elu", kernel_initializer="he_normal"),
keras.layers.Dropout(rate=0.2),
keras.layers.Dense(10,activation="softmax")

])

```

If you observed that the model is overfitting, you can increase the dropout rate. Conversely, you should try decreasing the dropout rate if the model underfits the training set. It can also help to increase the dropout rate for large layers, and reduce it for small ones. Moreover, many state-of-the-art architectures only use dropout after the last hidden layer, so you may want to try this if full dropout is too strong. Dropout does tend to significantly slow down convergence, but it usually result in a much better model when tuned properly. So, it is in generally well worth the extra time and effort.

observation Dropout is only active during training, made sure to evaluate the training loss without the dropout (after training).

MC Dropout

Max-Norm Regularization.

15.5. Summary

Lasso, Ridge, slow learning λ	tunning parameter
dropout learning	remove fraction of unit per layer and scaler weight
early stopping, epoch, batch size	details of stochastic gradient descent
number of layer, and unit fro layer	

15.6. bibliografy

2014 paper Nitish Srivastava (dropout)

Parte VI

Optimization (opt model).

Capítulo 16

First order methods.

16.1. Adagrad.

16.2. Adam.

Capítulo 17

Second order methods.

17.1. Newton's method.

17.2. Secant method.

Capítulo 18

Direct Methods.

18.1. Nelder-mead simplex method.

Capítulo 19

Stochastic methods.

19.1. Noise descent.

19.2. Cross-Entropy method.

Parte VII

Regularization (overfitting).

Parte VIII

Metrics (evaluation)

Capítulo 20

Regression

Capítulo 21

Classification

- Confusion matrix
- ROC.
- PR
- Gain and lift chart.

21.1. Matrix confusion

Fundamentally, the assessment process attempts to learn which model produces the most accurate and useful predictions.

When 1s are rare, the ratio of false positives to all predicted positives can be high, leading to the unintuitive situation in which a predicted 1 is most likely a 0.¹

In inbalance classes the rare class is usually the class of more interest and is typically designed 1, in contrast to more prevalent 0s. The 1s are the more important case, in the sense that misclassifying them as 0s is costlier than misclassifying 0s as 1s (exp legitimate insurance claims 0, fraudulent 1). The most accurate classification model may be one that simply classifies everything as a 0.

¹Default desicion point or cutoff is 0.5

		Predicted Class		
		Positive	Negative	
Actual Class	Positive	True Positive (TP)	False Negative (FN) Type II Error	Sensitivity $\frac{TP}{(TP + FN)}$
	Negative	False Positive (FP) Type I Error	True Negative (TN)	Specificity $\frac{TN}{(TN + FP)}$
		Precision $\frac{TP}{(TP + FP)}$	Negative Predictive Value $\frac{TN}{(TN + FN)}$	Accuracy $\frac{TP + TN}{(TP + TN + FP + FN)}$

Figura 21.1: Matrix cunfusion, sensitivity is recall.

21.2. ROC curve.

The Geometric Mean or G-Mean is a metric for imbalanced classification that, if optimized, will seek a balance between the sensitivity and the specificity.

$$G - mean = \sqrt{Sensitivity \cdot Specificity} \quad (21.1)$$

Given that we have already calculated the Sensitivity (TPR) and the complement to the Specificity when we calculated the ROC Curve, we can calculate the G-Mean for each threshold directly.

It turns out there is a much faster way to get the same result, called the Youden's J statistic.

$$J = Sensitivity + Specificity - 1 \quad (21.2)$$

We can then choose the threshold with the largest J statistic value





Figura 21.2: If ROC (sensity-recall) hugs the upperleft corner it will identify lots of 1s without misclassifying lots of 0s as 1s. Prevalence is $y=1 / \text{total}$. false positive rate is $1 - \text{specifity}$.



21.3. PR curve

PR curves are specially useful in evaluating data with highly unbalanced out-comes.




Figura 21.3:

Fuentes A Gentle Introduction to Threshold-Moving for Imbalanced Classification

If we are interested in a threshold that results in the best balance of precision and recall, then this is the same as optimizing the F-measure that summarizes the harmonic mean of both measures.


$$F - Measure = \frac{(2 \cdot Precision \cdot Recall)}{(Precision + Recall)} \quad (21.3)$$

first calculates the F-measure for each threshold, then locates the score and threshold with the largest value.

21.4. Gain chart and lift chart.

Gain and Lift charts are two approaches used while solving classifications problems with imbalanced data sets. The gain chart and lift chart is the measures in logistic regression that will help organizations to understand the benefits of using that model. Cumulative Gain Chart

$$Gain = \frac{CumulativeNumberOfPositiveObservationsUptodecibels}{TotalNumberOfPositiveObservationInData} \quad (21.4)$$



it says how much population we should sample to get the desired sensitivity of our classifier i.e. if we want to direct 40 % of potential repliers to our targeting campaign, we should select 20


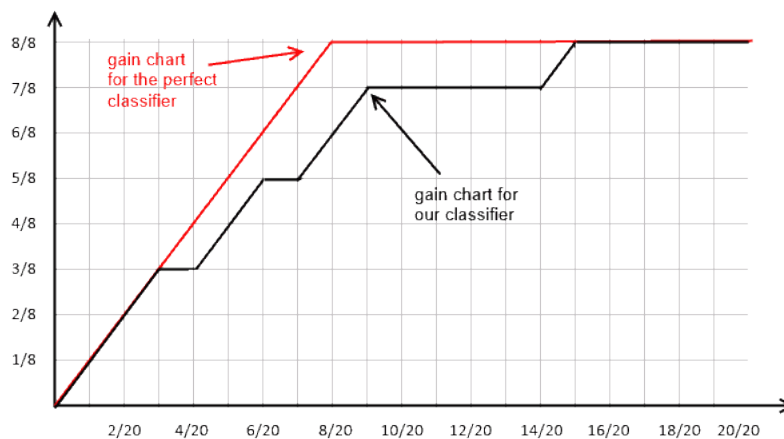


Figura 21.4: we can easily see if a classifier overfits on the test set, but underperforms on the testing

$$lift = \frac{Cumulative\ number\ of\ positive\ observations\ upto\ decile\ using\ ML\ model}{Cumulative\ Number\ of\ Positive\ Observations\ upto\ Decile\ using\ Random\ Model.} \quad (21.5)$$

falta analizar el lift, el libro donde encuentre es el data mining for eider franck, revisar luego, usare lo mencionado en ROC, PR y el tratamiendo de datos no balanceados por mientras, el lift es para saber los costos, vere mas adelante.

21.5. K-S Chart

21.6. bibliografia.

A Gentle Introduction to Threshold-Moving for Imbalanced Classification
 Cumulative Gain Chart
 Practical statistical for data science, peter bruce and andrew bruce

Parte IX

Python.

Capítulo 22

sklearn.preprocessing

Preprocessing and Normalization

The `sklearn.preprocessing` module includes scaling, centering, normalization, binarization methods.

- `sklearn.preprocessing.StandardScaler()`.

$$\frac{x_i - \text{mean}(x)}{\text{sd}(x)} \quad (22.1)$$

- `sklearn.preprocessing.MinMaxScaler()`.

$$\frac{x_i - \min(x)}{\max(x) - \min(x)} \quad (22.2)$$

- `sklearn.preprocessing.RobustScaler()`.

$$\frac{x_i - Q_1(x)}{Q_3(x) - Q_1(x)} \quad (22.3)$$




Figura 22.1: from Geeksfor geeks.

Capítulo 23

`sklearn.pipeline`

The `sklearn.pipeline` module implements utilities to build a composite estimator, as a chain of transforms and estimators.

Capítulo 24

`sklearn.ensemble`

The `sklearn.ensemble` module includes ensemble-based methods for classification, regression and anomaly detection.

24.1. `sklearn.ensemble.RandomForestClassifier`

Capítulo 25

`sklearn.model_selection`

See the Cross-validation: evaluating estimator performance, Tuning the hyper-parameters of an estimator and Learning curve sections for further details.

- `sklearn.model_selection.train_test_split`

Capítulo 26

sklearn.metrics

The `sklearn.metrics` module includes score functions, performance metrics and pairwise metrics and distance computations.

- `metrics.classification_report`
- `metrics.confusion_matrix`
- `metrics.roc_curve`
- `metrics.roc_auc_score`

Capítulo 27

`tf.keras.model`

Model groups layers into an object with training and inference features.

Capítulo 28

`ft.keras.layer`

Parte X

R.

Capítulo 29

ggplot

Capítulo 30

dplyr

Capítulo 31

tidyr

Parte XI

SQL.

Parte XII

Julia.

Parte XIII

C++.

Bibliografía

- [1] T. Hastie, R. Tibshirani, and J. Friedman, *The Elements of Statistical Learning: Data Mining, Inference, and Prediction*. Springer series in statistics, Springer, 2009.
- [2] G. James, D. Witten, T. Hastie, and R. Tibshirani, *An introduction to statistical learning*, vol. 112. Springer, 2013.